



# Università di Genova

## High Performance Computing Report OpenMP

Andrea Cattarinich 5137057

November 27, 2025

### Contents

|          |                                                              |           |
|----------|--------------------------------------------------------------|-----------|
| <b>1</b> | <b>Architectures</b>                                         | <b>2</b>  |
| <b>2</b> | <b>Hotspot Analysis</b>                                      | <b>3</b>  |
| 2.1      | Code Description . . . . .                                   | 3         |
| 2.2      | Project Structure . . . . .                                  | 4         |
| 2.3      | Sequential and Parallel Sections . . . . .                   | 5         |
| 2.4      | Compilation . . . . .                                        | 5         |
| 2.5      | Hotspot Identification . . . . .                             | 6         |
| 2.6      | Conclusions . . . . .                                        | 6         |
| <b>3</b> | <b>Amdahl's Law</b>                                          | <b>8</b>  |
| 3.1      | Performance Metrics . . . . .                                | 8         |
| 3.2      | The Law . . . . .                                            | 8         |
|          | Ideal case . . . . .                                         | 8         |
|          | Theoretical limit case . . . . .                             | 8         |
|          | Real case . . . . .                                          | 9         |
| 3.3      | Results . . . . .                                            | 9         |
| <b>4</b> | <b>Vectorization</b>                                         | <b>11</b> |
| 4.1      | Compilation Experiments through ICX . . . . .                | 11        |
| 4.2      | Conclusion . . . . .                                         | 13        |
| <b>5</b> | <b>Parallelization</b>                                       | <b>14</b> |
| 5.1      | Parallelization over the outer loop . . . . .                | 14        |
| 5.2      | Parallelization over the inner loop with reduction . . . . . | 14        |
| 5.3      | Results . . . . .                                            | 15        |
| <b>6</b> | <b>Conclusions</b>                                           | <b>16</b> |

# 1 Architectures

For this project I have used two workstation, my personal one and the workstation provided by the University.

## HP Envy 17-ae103nl (Laptop)

- **CPU:** Intel Core i7-6500U @ 2.50 GHz
- **Architecture:** x86\_64
- **Physical cores:** 2
- **Logical threads:** 4 (Hyper-Threading)

## University Workstation (Room 210)

- **CPU:** Intel Core i9-12900K (12th Gen)
- **Architecture:** x86\_64
- **CPU(s):** 24
- **Thread(s) per core:** 2
- **Core(s) per socket:** 16
- **Socket(s):** 1
- **CPU max MHz:** 4.02 GHz
- **Profiling Tools:** Intel VTune Profiler, Intel Advisor

To use the Intel oneAPI tools, the environment was loaded with:

```
source /opt/intel/oneapi/setvars.sh
```

## 2 Hotspot Analysis

The Fourier Transform is a fundamental mathematical tool in digital signal processing. It converts a sequence of values from the **time domain** into a sequence of values in the **frequency domain**. This transformation makes it possible to identify which frequencies are present in a signal and with what intensity.

The continuous-time Fourier Transform is defined as:

$$X(\omega) = \int_{-\infty}^{+\infty} x(t) e^{-j\omega t} dt$$

The example shown in Figure 1 illustrates the practical and optical interpretation of the Fourier Transform of a rectangular signal, showing how a simple time-domain shape corresponds to a characteristic frequency-domain spectrum.




Figure 1: Fourier Transform of a rectangular signal.

### 2.1 Code Description

The algorithm provided by the course implements both the Discrete Fourier Transform (DFT) and its inverse (IDFT) in C. It computes the transform of a signal and verifies the correctness by applying the IDFT, following the discrete-time Fourier Transform definition:

$$X[k] = \sum_{n=0}^{N-1} x[n] e^{-j2\pi kn/N}, \quad k = 0, 1, \dots, N-1$$

The DFT is implemented following the standard quadratic-complexity  $O(N^2)$  formulation, with two nested loops over the frequency index  $k$  and the time index  $n$ . Each output element can be computed independently, which makes the algorithm naturally parallelisable.

```

1 int main() {
2     int N = 100000;
3     double *xr, *xi, *Xr, *Xi;
4
5     fillInput(xr, xi, N);           // initialize input signal
6     zero(Xr, Xi, N);               // clear output buffers
7
8     DFT(+1, xr, xi, Xr, Xi, N);     // compute DFT
9     DFT(-1, Xr, Xi, xr, xi, N);     // compute IDFT
10
11     checkResults(xr, xi, Xr, Xi, N); // verify correctness
12
13     return 0;
14 }
15
16 void DFT(int idft, double* xr, ..., double* Xi, int N) {
17     for (int k = 0; k < N; k++)
18         for (int n = 0; n < N; n++) {
19             double a = 2*M_PI*n*k/N;
20             Xr[k] += xr[n]*cos(a) + idft*xi[n]*sin(a);
21             Xi[k] += -idft*xr[n]*sin(a) + xi[n]*cos(a);
22         }
23
24     if (idft == -1)
25         for (int n = 0; n < N; n++) {
26             Xr[n] /= N;
27             Xi[n] /= N;
28         }
29 }

```

Listing 1: Pseudo code in C

## 2.2 Project Structure

The `homework1` project is organized into four main directories:

- **build**: contains the compiled executables.
- **reports**: stores the reports generated with the `-qopt-report` flag.
- **src**: includes the source code files.
- **metrics**: scripts that run the executables in **build** and save performance metrics.

The first step to improving the performance of a program is profiling. The goal of this analysis is to identify the parts of the program that consume the most execution time. To achieve this, I considered a data size resulting in a sequential execution time of  $\sim 30$  seconds, by setting  $N = 60'000$ .

From this point onward, each executable is empirically measured over 10 runs to obtain reliable results.

## 2.3 Sequential and Parallel Sections

In order to apply Amdahl's Law (discussed in more detail in the Chapter 3), we must identify the sequential and parallel parts of the code. I have empirically measured these sections by separating the serial and parallel portions. More details of how I obtained these results can be found in the code `sequential_parallel_code_identification.c`.

The results are the following:

| Quantity                                                                      | Value      |
|-------------------------------------------------------------------------------|------------|
| Total execution time $T_{\text{tot}}$                                         | 32.48231 s |
| Time of the parallelizable section $T_{\text{par}}$                           | 32.48021 s |
| Time of the serial section $T_{\text{ser}} = T_{\text{tot}} - T_{\text{par}}$ | 0.00209 s  |
| Fraction of serial code $f_s = \frac{T_{\text{ser}}}{T_{\text{tot}}}$         | 0.006 %    |

These values will be used to estimate the theoretical speedup according to Amdahl's Law:

$$S(p) = \frac{1}{f_s + \frac{1-f_s}{p}} \Big|_{f_s=0.00006} = \frac{1}{0.00006 + \frac{1-0.00006}{p}} \approx p$$

**Note:** the smaller the serial fraction  $f_s$ , the higher the potential speedup achievable by parallelizing the main loop of the DFT.

**Note: in the real case, the speedup is not linear.** In particular, from the parallelization examples presented in Chapter 5, the speedup is approximately linear up to 8 threads; beyond this point, it deviates from linear behavior.

## 2.4 Compilation

Since the Intel(R) C++ Compiler Classic (ICC) is deprecated and was removed from product release in the second half of 2023. The Intel(R) oneAPI DPC++/C++ Compiler (ICX) was used to performe the analysis.

To identify the optimal sequential compilation strategy, I conduced testing across multiple Intel compiler variants (ICC and ICX), focusing on the vectorization report provided by the flag `-qopt-report`.

The program was compiled using the Intel compiler ICX with the following command (more details in Chapter 4):

```
icx -O2 -vec -g \
    -fopenmp \
    -qopt-report=max \
    -qopt-report-file=reports/omp_homework.optrpt \
    -o build/omp_homework \
    src/omp_homework.c
```

The compilation flags have the following meaning:

- `-O2`: optimizations for better performance;

- `-fopenmp`: activates compilation for OpenMP directives;
- `-vec`: activates automatic vectorization of loops using SIMD instructions;
- `-g`: includes debugging information in the executable;
- `-qopt-report=max`: generates a detailed optimization report;
- `-qopt-report-file=...`: specifies the file where the optimization report is saved;
- `-o build/omp_homework`: sets the name and location of the compiled executable;
- `src/omp_homework.c`: specifies the source file to compile.

## 2.5 Hotspot Identification

The measurements were performed using two methods: **Manually** (using `omp_get_wtime()`) and **Profiler** (using Intel VTune)

Table 1: Manually

| Function       | Time (s) | Percentage (%) |
|----------------|----------|----------------|
| fillInput      | 0.0020   | 0.0061         |
| setOutputZero1 | 0.0000   | 0.0000         |
| setOutputZero2 | 0.0000   | 0.0000         |
| DFT            | 16.2665  | 50.0319        |
| IDFT           | 16.2435  | 49.9613        |
| Total          | 32.5122  | 100.0000       |

Table 2: Intel VTune Profiler

| Function Stack         | CPU Time: Total | Percentage (%) |
|------------------------|-----------------|----------------|
| Total                  | 32.4300         | 100.0000       |
| ._start                | 32.4300         | 100.0000       |
| ._libc_start_main_impl | 32.4300         | 100.0000       |
| - main                 | 32.4300         | 100.0000       |
| — DFT                  | 16.2420         | 50.0833        |
| — __svml_sincos2       | 9.3781          | 28.9180        |
| — __svml_sincos2_l9    | 2.7880          | 8.5969         |
| — DFT                  | 15.8800         | 48.9670        |
| — __svml_sincos2       | 9.7499          | 30.0645        |
| — __svml_sincos2_l9    | 2.4501          | 7.5549         |
| — fillInput            | 0.3080          | 0.9497         |

### Top Hotspots

This section lists the most active functions in your application. Optimizing these hotspot functions typically results in improving overall application performance.

| Function                       | Module                  | CPU Time ☺ | % of CPU Time ☺ |
|--------------------------------|-------------------------|------------|-----------------|
| <a href="#">__svml_cos2_l9</a> | hotspots_identification | 18.344s    | 56.6%           |
| <a href="#">DFT</a>            | hotspots_identification | 4.076s     | 12.6%           |
| <a href="#">DFT</a>            | hotspots_identification | 3.680s     | 11.3%           |
| <a href="#">__svml_sin2_l9</a> | hotspots_identification | 3.386s     | 10.4%           |
| <a href="#">__svml_sin2</a>    | hotspots_identification | 1.852s     | 5.7%            |
| [Others]                       | N/A*                    | 1.092s     | 3.4%            |

\*N/A is applied to non-summable metrics.

## 2.6 Conclusions

The main hotspot is the `DFT()` function and its inverse `IDFT()`, which together account for over 99,05% of the total execution time. All other functions have negligible impact. This indicates that **parallelization and vectorization efforts should focus on the**

double loop inside DFT().

The main hotspot is the loop contained in lines 72-80 (source: src/omp\_homework.c).

The dominant operations contributing to the algorithm's computational intensity are the `sin()` and `cos()` trigonometric functions, which account for the majority of the execution time, as shown in the following code.

```
// DFT/IDFT routine
// idft: 1 direct DFT, -1 inverse IDFT (Inverse DFT)
int DFT(int idft, double* xr, double* xi, double* Xr_o, double* Xi_o, int N){
    int k, n;
    for (k=0 ; k<N ; k++)
    {
        for (n=0 ; n<N ; n++) {
            // Real part of X[k]
            Xr_o[k] += xr[n] * cos(n * k * PI2 / N) + idft*xi[n]*sin(n * k * PI2 / N);
            // Imaginary part of X[k]
            Xi_o[k] += -idft*xr[n] * sin(n * k * PI2 / N) + xi[n] * cos(n * k * PI2 / N);
        }
    }

    // normalize if you are doing IDFT
    if (idft==-1){
        for (n=0 ; n<N ; n++){
            Xr_o[n] /=N;
            Xi_o[n] /=N;
        }
    }

    return 1;
}
```

### 3 Amdahl's Law

This section discusses Amdahl's Law, a fundamental principle in parallel computing that helps predict the theoretical speedup in latency when using multiple processors.

#### 3.1 Performance Metrics

Considering  $T_p$  the time to solve a problem using  $p$  processors, we define:

- **Speedup:**  $S(p) = \frac{T_1}{T_p}$ ,  $0 < S(p) \leq p$
- **Efficiency:**  $E(p) = \frac{S(p)}{p}$ ,  $0 < E(p) \leq 1$

#### 3.2 The Law

Amdahl's Law can be expressed as a classical formula:

$$S(p) = \frac{1}{f_s + \frac{(1-f_s)}{p}}$$

where:

- $S(p)$ : is the theoretical speedup
- $f_s$ : is the proportion of sequential code (non-parallelizable)
- $1 - f_s$ : is the proportion of parallelizable code
- $p$ : is the number of processors

**Ideal case** ( $f_s = 0$ )

Implies that all code is parallelizable. The speedup becomes:

$$S(p) = \frac{1}{\frac{1}{p}} = p$$

Maximum speedup equals the number of processors. Each processor contributes fully to the speedup.

**Theoretical limit case** ( $p \rightarrow \infty$ )

By increasing the number of processors to infinity, the speedup reaches its theoretical limit:

$$\lim_{p \rightarrow \infty} S(p) = \frac{1}{f_s}$$

Even though the number of processors increases indefinitely, the speedup is limited by the sequential portion of the code, encountering a bottleneck.



**Real case** ( $f_s = \alpha p$ )

Where  $\alpha$  is a constant  $\alpha \leq 0.1$  representing the overhead per processor. The speedup becomes:

$$S(p) = \frac{1}{\alpha p + \frac{1-\alpha p}{p}} = \frac{p}{1 - \alpha p + \alpha p^2}$$

There is no free lunch: as we add more processors, the overhead increases linearly, impacting the overall speedup.

**For  $\alpha$  small:** when  $\alpha$  is small we do not have a significant overhead.

$$S(p) \approx p$$

**For  $\alpha$  large:** when  $\alpha$  is large, the overhead dominates. The more we add processors, the less benefit we get!


$$S(p) \approx \frac{p}{\alpha p^2} \approx \frac{1}{\alpha p}$$

### 3.3 Results

The following examples have been obtained using Matlab scripts available in the repository (`classic_vs_overhead.m` and `parallel_comparison.m`).


#### Classical vs Overhead

- The proportion of the parallelizable is 95% of code.  
( $f_s = 0.05$ ,  $f_p = 1 - f_s = 0.95$ )
- Overhead per processor  
 $\alpha = 0.01$



## Parallel speedup comparison

It is useful for understanding the limits of parallelization according to Amdahl's law: it highlights that simply increasing the number of processors is not enough if the sequential portion of the code is significant



## 4 Vectorization

The vectorization of the baseline code (`src/omp_homework.c`) was analyzed using Intel Advisor, focusing on performance hotspots and Roofline Analysis. The workflow included adding the executable and source folders, selecting **Roofline Analysis**, checking **Trip Counts**, and running the analysis.

### 4.1 Compilation Experiments through ICX

#### 1. Baseline Compilation (no optimization)

Compilation flags: `-g -fopenmp`

Hotspots are memory-bound (DRAM), with no vectorization applied. Roofline shows performance limited by the Scalar Add Peak ( $\sim 7.97$  GFLOPS).




Figure 2: Roofline Analysis - baseline compilation.

#### 2. Basic Optimization

Compilation flags: `-O0 -g -fopenmp`

Arithmetic intensity improves from 0.054 to 0.213 FLOP/byte. DRAM-bound, but performance increases relative to baseline. Hotspots correspond to trigonometric function calls.




Figure 3: Roofline Analysis - basic optimization.

### 3. Best-Optimized Compilation

Compilation flags: `-O2 -vec -g -fopenmp -xHost`

SIMD vectorization applied to critical loops. Performance rises from 10.1 GFLOPS to 34.5 GFLOPS while arithmetic intensity remains similar.




Figure 4: Roofline Analysis - best-optimized compilation.

### 4. Parallelization with Optimization

OpenMP parallelization over the outer loop (discussed in Chapter 5) further boosts performance. Roofline Analysis indicates:

| Function | Arithmetic Intensity | Performance (GFLOPS) |
|----------|----------------------|----------------------|
| cos      | 0.213                | 214,817              |
| sin      | 0.232                | 418,012              |

Table 3: Performance of trigonometric functions with parallelization.




Figure 5: Roofline Analysis - parallelized and optimized code.

## 4.2 Conclusion

Roofline Analysis and Intel Advisor guided the optimization process:

- Baseline code is memory-bound with no vectorization.
- Compiler optimizations and architecture-specific flags increase performance and apply SIMD efficiently.
- Combining vectorization with OpenMP parallelization allows computationally intensive functions to approach hardware limits.

## 5 Parallelization

### 5.1 Parallelization over the outer loop

In the first approach (`src/omp1.c`), the outer loop over the frequency index  $k$  is parallelized.

```
1 #pragma omp parallel for private(n)
2 for (k = 0; k < N; k++) {
3     double Xr_temp = 0.0;
4     double Xi_temp = 0.0;
5     for (n = 0; n < N; n++) {
6         Xr_temp += xr[n] * cos(n * k * PI2 / N) + idft * xi[n] *
7             sin(n * k * PI2 / N);
8         Xi_temp += -idft * xr[n] * sin(n * k * PI2 / N) + xi[n] *
9             cos(n * k * PI2 / N);
10    }
11    Xr_o[k] = Xr_temp;
12    Xi_o[k] = Xi_temp;
13 }
```

Here, each thread computes a subset of the output frequencies independently. The temporary variables `Xr_temp` and `Xi_temp` are private to each thread, ensuring there is no race condition when updating the output arrays.

### 5.2 Parallelization over the inner loop with reduction

In the second approach, the inner loop over the time index  $n$  is parallelized using a reduction clause:

```
1 for (k = 0; k < N; k++) {
2     double Xr_temp = 0.0;
3     double Xi_temp = 0.0;
4     #pragma omp parallel for reduction(+:Xr_temp, Xi_temp)
5     for (n = 0; n < N; n++) {
6         Xr_temp += xr[n] * cos(n * k * PI2 / N) + idft * xi[n] *
7             sin(n * k * PI2 / N);
8         Xi_temp += -idft * xr[n] * sin(n * k * PI2 / N) + xi[n] *
9             cos(n * k * PI2 / N);
10    }
11    Xr_o[k] = Xr_temp;
12    Xi_o[k] = Xi_temp;
13 }
```

Here, the computation of each output frequency  $X[k]$  is parallelized across the time samples  $n$ . The `reduction` clause ensures that each thread accumulates contributions to `Xr_temp` and `Xi_temp` safely, avoiding race conditions. This method can provide higher parallel efficiency for large  $N$ , but incurs some overhead from the reduction operation, as shown in the following figures.

## 5.3 Results


Parallelized for loop with OpenMP using `private(n)` to make each thread independent.

```


1 #pragma omp parallel for private(n)
2 for (k=0 ; k<N ; k++)
3 {
4     double Xr_temp = 0.0;
5     double Xi_temp = 0.0;
6
7     for (n=0 ; n<N ; n++) {
8         Xr_temp += ...
9         Xi_temp += ...
10    }
11
12    Xr_o[k] = Xr_temp;
13    Xi_o[k] = Xi_temp;
14 }

```


Execution time



Speedup



Efficiency




Parallelized for loop with OpenMP using `reduction` to combine the results from all threads.

```


1 int k, n;
2 for (k=0 ; k<N ; k++)
3 {
4     double Xr_temp = 0.0;
5     double Xi_temp = 0.0;
6
7 #pragma omp parallel for reduction(+:
8     Xr_temp, Xi_temp)
9     for (n=0 ; n<N ; n++) {
10         Xr_temp += ...
11         Xi_temp += ...
12    }
13
14    Xr_o[k] = Xr_temp;
15    Xi_o[k] = Xi_temp;
16 }

```


Execution time



Speedup



Efficiency



## Attempted Optimization

**Precomputation of Trigonometric Functions:** in an additional attempt to optimize the DFT computation, I modified the code as in the source `src/omp3.c`. The idea was to reduce the repeated evaluation of trigonometric functions by computing the sine and cosine values once per iteration:

```
1 #pragma omp parallel for private(n)
2 for (k=0; k<N; k++) {
3     double Xr_temp = 0.0;
4     double Xi_temp = 0.0;
5
6 #pragma omp simd reduction(+:Xr_temp, Xi_temp)
7     for (n=0; n<N; n++) {
8         double angle = n * k * PI2 / N;
9         double c = cos(angle);
10        double s = sin(angle);
11
12        Xr_temp += xr[n]*c + idft*xi[n]*s;
13        Xi_temp += -idft*xr[n]*s + xi[n]*c;
14    }
15
16    Xr_o[k] = Xr_temp;
17    Xi_o[k] = Xi_temp;
18 }
```

The motivation behind this modification was that, theoretically, computing the trigonometric functions only once per iteration should reduce the overall workload and improve performance.

However, the experiments did not yield significant improvements. Performance remained the same, or in some cases slightly worse. This is likely due to the compiler automatically applying this optimization by default.

## 6 Conclusions

Both parallelization approaches exhibit near-linear scaling with an efficiency of approximately 100% up to 8 cores. The first parallelization strategy demonstrated the best performance, achieving a maximum speedup of 7.87 on 8 cores. Beyond 24 cores, additional cores provided no significant improvement.